

Capstone-Heart-RBM-Final

April 7, 2023

```
[1]: # Capstone Project - Heart Disease Prediction
# Dataset - data.xlsx in folder '1582800613_project3datadictionary' as
↳downloaded from website
```

```
[2]: # Initial preparation to work on Python
# Import packages pandas, numpy and os

import pandas as pd
import numpy as np
import os
```

```
[3]: # Get current working directory

os.getcwd()

# Current working directory will be used for this project
```

```
[3]: '/home/labsuser/Capstone'
```

```
[4]: os.chdir('/home/labsuser/Capstone')
```

```
[5]: os.getcwd()
```

```
[5]: '/home/labsuser/Capstone'
```

```
[6]: # Read the dataset

heart = pd.read_excel('data.xlsx')
```

```
[7]: # Output the dataset

heart

# dataset has 303 rows and 14 columns
```

```
[7]:      age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  \
0      63    1   3      145    233    1         0      150     0       2.3
```

1	37	1	2	130	250	0	1	187	0	3.5
2	41	0	1	130	204	0	0	172	0	1.4
3	56	1	1	120	236	0	1	178	0	0.8
4	57	0	0	120	354	0	1	163	1	0.6
..
298	57	0	0	140	241	0	1	123	1	0.2
299	45	1	3	110	264	0	1	132	0	1.2
300	68	1	0	144	193	1	1	141	0	3.4
301	57	1	0	130	131	0	1	115	1	1.2
302	57	0	1	130	236	0	0	174	0	0.0

	slope	ca	thal	target
0	0	0	1	1
1	0	0	2	1
2	2	0	2	1
3	2	0	2	1
4	2	0	2	1
..
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

[303 rows x 14 columns]

```
[8]: # Initial check of the data
```

```
heart.head()
```

```
[8]:   age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope  \
0   63   1   3     145    233   1         0     150     0      2.3     0
1   37   1   2     130    204   0         0     172     0      1.4     2
2   41   0   1     130    204   0         1     178     0      0.8     2
3   56   1   1     120    236   0         1     163     1      0.6     2
4   57   0   0     120    354   0         1     163     1      0.6     2
```

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
[9]: # Initial check of the data
```

```
heart.tail()
```

```
[9]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
298	57	0	0	140	241	0	1	123	1	0.2	
299	45	1	3	110	264	0	1	132	0	1.2	
300	68	1	0	144	193	1	1	141	0	3.4	
301	57	1	0	130	131	0	1	115	1	1.2	
302	57	0	1	130	236	0	0	174	0	0.0	

	slope	ca	thal	target
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

```
[10]: # Check for Null values or missing data

heart.isna().sum()

# There are no null values
# There are Zero values which are important to the dataset for analysis
# They need to be retained
```

```
[10]: age          0
sex          0
cp           0
trestbps     0
chol         0
fbs          0
restecg      0
thalach      0
exang        0
oldpeak      0
slope        0
ca           0
thal         0
target       0
dtype: int64
```

```
[11]: # Check for data types

heart.dtypes

# Data is all integer except for the oldpeak which is float64
# Data to be retained in current format
```

```
[11]: age          int64
sex          int64
```

```

cp                int64
trestbps          int64
chol              int64
fbs               int64
restecg           int64
thalach           int64
exang             int64
oldpeak           float64
slope             int64
ca                int64
thal              int64
target            int64
dtype: object

```

[12]: *# Additional check for data integrity*

```

heart.info()

# 303 non-null values in dataset with 14 columns - all integer except oldpeak
→ is float64
# 13 independent variables and 1 dependent variable - 'target' (which needs to
→ be predicted)

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps     303 non-null   int64
4   chol        303 non-null   int64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   int64
8   exang       303 non-null   int64
9   oldpeak     303 non-null   float64
10  slope       303 non-null   int64
11  ca          303 non-null   int64
12  thal        303 non-null   int64
13  target      303 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB

```

[13]: *# Dataset structure - 13 variables as under:*

```

# Categorical values - 7 - sex, cp, restecg, exang, slope, ca and thal

```

```
# Boolean values - 1 - fbs
# Numeric values - 5 - age, trestbps, chol, thalach, oldpeak
```

```
[14]: # Check the value counts for the target variable
```

```
heart['target'].value_counts()

# Data indicates heart disease risk higher for 165 and low for 138
# There is good balance between positive and negative CVD cases
```

```
[14]: 1    165
      0    138
      Name: target, dtype: int64
```

```
[15]: # Check the number of males and females in the dataset
```

```
heart['sex'].value_counts()

# There are 207 males and 96 females
```

```
[15]: 1    207
      0    96
      Name: sex, dtype: int64
```

```
[16]: # Import the plot packages
```

```
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[17]: # Check the statistical output of all the variables
```

```
heart.describe()
```

```
[17]:
```

	age	sex	cp	trestbps	chol	fbs \
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000

	restecg	thalach	exang	oldpeak	slope	ca \
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	0.528053	149.646865	0.326733	1.039604	1.399340	0.729373
std	0.525860	22.905161	0.469794	1.161075	0.616226	1.022606

min	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	133.500000	0.000000	0.000000	1.000000	0.000000
50%	1.000000	153.000000	0.000000	0.800000	1.000000	0.000000
75%	1.000000	166.000000	1.000000	1.600000	2.000000	1.000000
max	2.000000	202.000000	1.000000	6.200000	2.000000	4.000000

	thal	target
count	303.000000	303.000000
mean	2.313531	0.544554
std	0.612277	0.498835
min	0.000000	0.000000
25%	2.000000	0.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	3.000000	1.000000

```
[18]: # Check the correlation between the variables to understand the relationship
      ↪ between all the variables
```

```
heart.corr()
```

```
[18]:
```

	age	sex	cp	trestbps	chol	fbs	\
age	1.000000	-0.098447	-0.068653	0.279351	0.213678	0.121308	
sex	-0.098447	1.000000	-0.049353	-0.056769	-0.197912	0.045032	
cp	-0.068653	-0.049353	1.000000	0.047608	-0.076904	0.094444	
trestbps	0.279351	-0.056769	0.047608	1.000000	0.123174	0.177531	
chol	0.213678	-0.197912	-0.076904	0.123174	1.000000	0.013294	
fbs	0.121308	0.045032	0.094444	0.177531	0.013294	1.000000	
restecg	-0.116211	-0.058196	0.044421	-0.114103	-0.151040	-0.084189	
thalach	-0.398522	-0.044020	0.295762	-0.046698	-0.009940	-0.008567	
exang	0.096801	0.141664	-0.394280	0.067616	0.067023	0.025665	
oldpeak	0.210013	0.096093	-0.149230	0.193216	0.053952	0.005747	
slope	-0.168814	-0.030711	0.119717	-0.121475	-0.004038	-0.059894	
ca	0.276326	0.118261	-0.181053	0.101389	0.070511	0.137979	
thal	0.068001	0.210041	-0.161736	0.062210	0.098803	-0.032019	
target	-0.225439	-0.280937	0.433798	-0.144931	-0.085239	-0.028046	

	restecg	thalach	exang	oldpeak	slope	ca	\
age	-0.116211	-0.398522	0.096801	0.210013	-0.168814	0.276326	
sex	-0.058196	-0.044020	0.141664	0.096093	-0.030711	0.118261	
cp	0.044421	0.295762	-0.394280	-0.149230	0.119717	-0.181053	
trestbps	-0.114103	-0.046698	0.067616	0.193216	-0.121475	0.101389	
chol	-0.151040	-0.009940	0.067023	0.053952	-0.004038	0.070511	
fbs	-0.084189	-0.008567	0.025665	0.005747	-0.059894	0.137979	
restecg	1.000000	0.044123	-0.070733	-0.058770	0.093045	-0.072042	
thalach	0.044123	1.000000	-0.378812	-0.344187	0.386784	-0.213177	
exang	-0.070733	-0.378812	1.000000	0.288223	-0.257748	0.115739	


oldpeak	-0.058770	-0.344187	0.288223	1.000000	-0.577537	0.222682
slope	0.093045	0.386784	-0.257748	-0.577537	1.000000	-0.080155
ca	-0.072042	-0.213177	0.115739	0.222682	-0.080155	1.000000
thal	-0.011981	-0.096439	0.206754	0.210244	-0.104764	0.151832
target	0.137230	0.421741	-0.436757	-0.430696	0.345877	-0.391724

	thal	target
age	0.068001	-0.225439
sex	0.210041	-0.280937
cp	-0.161736	0.433798
trestbps	0.062210	-0.144931
chol	0.098803	-0.085239
fbs	-0.032019	-0.028046
restecg	-0.011981	0.137230
thalach	-0.096439	0.421741
exang	0.206754	-0.436757
oldpeak	0.210244	-0.430696
slope	-0.104764	0.345877
ca	0.151832	-0.391724
thal	1.000000	-0.344029
target	-0.344029	1.000000

[19]: *# For better visualization of the correlation, convert the above into a heat map*

```
plt.figure(figsize=(20,12))
sns.set_context('notebook',font_scale = 1.3)
sns.heatmap(heart.corr(),annot=True,linewidth =2)
plt.tight_layout()
```

```
# The following can be inferred from the heat map:
# 1. cp, restecg, thalach and slope have a high positive correlation with the
    ↪target variable
# 2. exang - exercise induced angina and oldpeak - ST depression induced by
    ↪exercise relative to rest -
# - have high negative correlation to target indicating more stress/blood
    ↪requirement by heart during exercise
# Other variables are negatively correlated to the target variable
# 3. There is also a higher positive correlation between variables as below:
#    a. thalach and cp
#    b. slope and thalach
# thalach - highset heart rate achieved therefore needs further analysis
```



```
[20]: # Carry out univariate analysis
# AGE
# Age distribution - min, max, median
# Age classification into young, middle and old ages
# Create various plots to analyse the data
# 1. Age distribution

min(heart.age)
```

[20]: 29

```
[21]: max(heart.age)
```

[21]: 77

```
[22]: heart.age.mean()
```

[22]: 54.366336633663366

```
[23]: # Minimum age in the dataset is 29
# Maximum age in dataset is 77
# Mean age in the dataset is 54.36
```



```
[24]: # Grouping data to create bar plot
```

```
age_freq = heart.groupby('age').agg({'age': 'count'})  
age_freq
```

```
[24]:
```


	age
age	
29	1
34	2
35	4
37	2
38	3
39	4
40	3
41	10
42	8
43	8
44	11
45	8
46	7
47	5
48	7
49	5
50	7
51	12
52	13
53	8
54	16
55	8
56	11
57	17
58	19
59	14
60	11
61	8
62	11
63	9
64	10
65	8
66	7
67	9
68	4
69	3
70	4
71	3
74	1
76	1

[25]: *# Create a bar plot for age distribution*

```
plt.figure(figsize = (20,10))
barplot = plt.bar(age_freq.index, age_freq.age, fc = 'green', ec = 'black',
    ↪width = 0.7)
plt.title('Age Distribution')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.ylim([0,22])
plt.xticks(age_freq.index)
plt.bar_label(barplot, labels = age_freq.age, label_type = 'edge', padding = 3)
```

[25]: [Text(0, 3, '1'),
Text(0, 3, '2'),
Text(0, 3, '4'),
Text(0, 3, '2'),
Text(0, 3, '3'),
Text(0, 3, '4'),
Text(0, 3, '3'),
Text(0, 3, '10'),
Text(0, 3, '8'),
Text(0, 3, '8'),
Text(0, 3, '11'),
Text(0, 3, '8'),
Text(0, 3, '7'),
Text(0, 3, '5'),
Text(0, 3, '7'),
Text(0, 3, '5'),
Text(0, 3, '7'),
Text(0, 3, '12'),
Text(0, 3, '13'),
Text(0, 3, '8'),
Text(0, 3, '16'),
Text(0, 3, '8'),
Text(0, 3, '11'),
Text(0, 3, '17'),
Text(0, 3, '19'),
Text(0, 3, '14'),
Text(0, 3, '11'),
Text(0, 3, '8'),
Text(0, 3, '11'),
Text(0, 3, '9'),
Text(0, 3, '10'),
Text(0, 3, '8'),
Text(0, 3, '7'),

```
Text(0, 3, '9'),
Text(0, 3, '4'),
Text(0, 3, '3'),
Text(0, 3, '4'),
Text(0, 3, '3'),
Text(0, 3, '1'),
Text(0, 3, '1'),
Text(0, 3, '1')]
```



```
[26]: # Categorizing the ages into 3 groups to understand the spread of data within
      ↪ the dataset
      # Categorize into 3 categories - Young (25-40 yrs), Middle (40-55 yrs) and Old
      ↪ (>55 yrs)
      # Create an additional column with the categorization

heart['Age_Category'] = np.where((heart.age >= 29) & (heart.age < 40), 'Young
      ↪ Age',
                                np.where((heart.age >= 40) & (heart.age <
      ↪ 55), 'Middle Age', 'Old Age'))
heart
```

```
[26]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63	1	3	145	233	1	0	150	0	2.3	
1	37	1	2	130	250	0	1	187	0	3.5	
2	41	0	1	130	204	0	0	172	0	1.4	
3	56	1	1	120	236	0	1	178	0	0.8	
4	57	0	0	120	354	0	1	163	1	0.6	

298	57	0	0	140	241	0	1	123	1	0.2
299	45	1	3	110	264	0	1	132	0	1.2
300	68	1	0	144	193	1	1	141	0	3.4
301	57	1	0	130	131	0	1	115	1	1.2
302	57	0	1	130	236	0	0	174	0	0.0

	slope	ca	thal	target	Age_Category
0	0	0	1	1	Old Age
1	0	0	2	1	Young Age
2	2	0	2	1	Middle Age
3	2	0	2	1	Old Age
4	2	0	2	1	Old Age

298	1	0	3	0	Old Age
299	1	0	3	0	Middle Age
300	1	2	3	0	Old Age
301	1	1	3	0	Old Age
302	1	1	2	0	Old Age

[303 rows x 15 columns]

```
[27]: # Aggregate the data with the new categorization

age_category = heart.groupby('Age_Category').agg({'Age_Category' : 'count'}).
    ↳ apply(lambda x: x.sort_values(ascending = True).head(3))
age_category
```

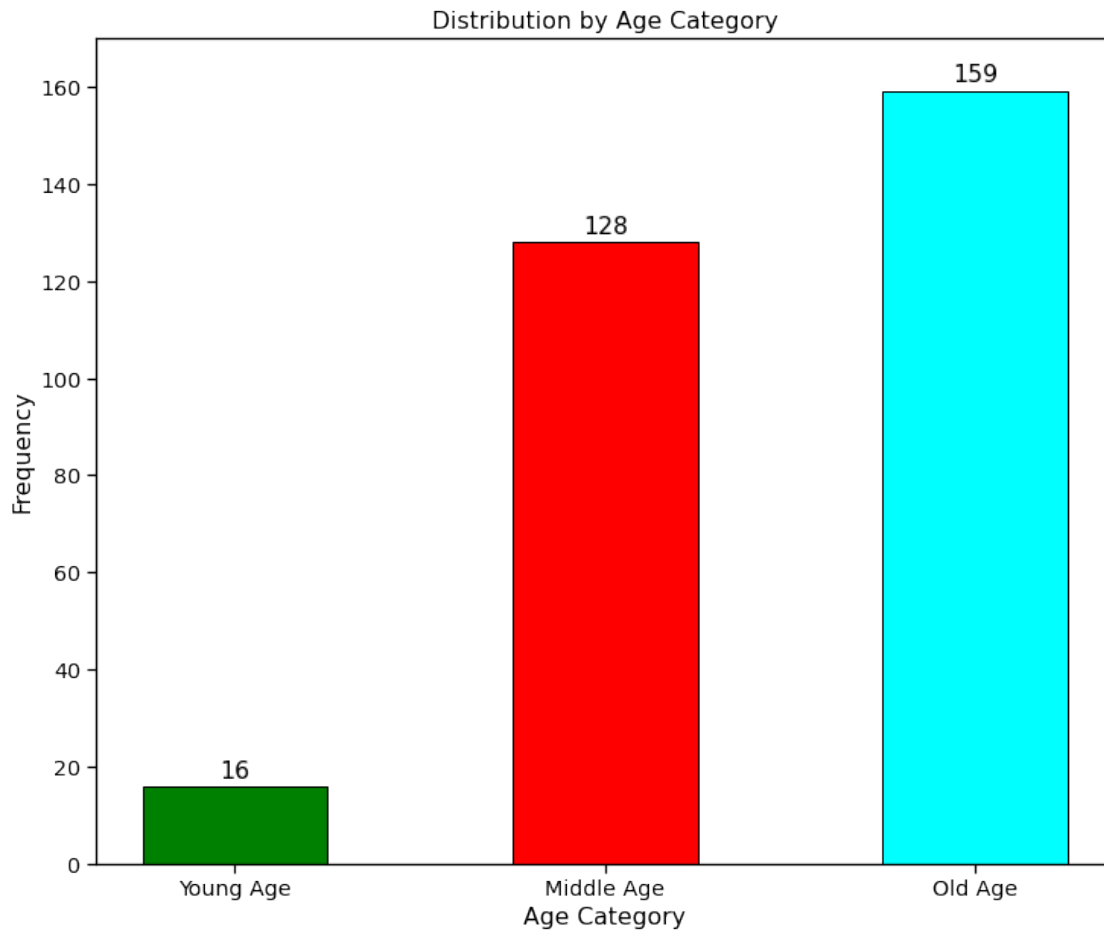
```
[27]:          Age_Category
Age_Category
Young Age          16
Middle Age         128
Old Age            159
```

```
[28]: # Plot a barplot to see the distribution of samples in the dataset by age_
    ↳ category

plt.figure(figsize = (12,10))
barplot = plt.bar(age_category.index, age_category.Age_Category,
                  color = ['green', 'red', 'cyan'], ec = 'black', width = 0.5)
plt.title('Distribution by Age Category')
plt.xlabel('Age Category')
plt.ylabel('Frequency')
plt.ylim([0,170])
plt.xticks(age_category.index)
plt.bar_label(barplot, labels = age_category.Age_Category, label_type = 'edge',
    ↳ padding = 3)
```

```
# It is seen that the dataset has maximum old age - 159, middle age - 128 and
↳ young age - 16
```

```
[28]: [Text(0, 3, '16'), Text(0, 3, '128'), Text(0, 3, '159')]
```



```
[29]: # Understand and check the relation between high risk CVD to the age group

age_group = heart.groupby('Age_Category')['target'].apply(lambda x: (x == 1).
↳sum()).reset_index(name = 'Total_CVD')
age_group
```

```
# In the Young age group, there are 12 cases with CVD
# In the Middle age group, there are 88 cases of CVD
# In the Old age group, there are 65 cases of CVD
```

```
[29]:   Age_Category  Total_CVD
0    Middle Age          88
```

1	Old Age	65
2	Young Age	12

```
[30]: # Create a new dataframe to display the overall statistics of CVD with respect
      ↪to age

age_grouptotal = pd.DataFrame({'Age_Category' : ['Young Age', 'Middle Age',
      ↪'Old Age'],
                               'Total_CVD' : [12, 88, 65], 'Total' : [16, 128,
      ↪159]})
age_grouptotal
```

```
[30]:  Age_Category  Total_CVD  Total
0    Young Age         12      16
1  Middle Age         88     128
2    Old Age          65     159
```


```
[31]: # Create a pie chart to understand the % distribution of the population

plt.figure(figsize = (10,10))
plt.pie(age_grouptotal.Total, labels = age_grouptotal.Age_Category, autopct =
      ↪'%0.1f%%')
plt.title('Distribution by Age Category')
plt.legend()

# Population distribution is as follows:
# 1. Young age - 5.3% of the total 303
# 2. Middle age - 42.2% of the total 303
# 3. Old age - 52.5% of the total 303
```

```
[31]: <matplotlib.legend.Legend at 0x7f4aadfd0f10>
```

Distribution by Age Category



```
[32]: # Add a % of the population column to the dataframe

age_grouptotal['Percent_Population'] = round((age_grouptotal['Total_CVD']/
                                              age_grouptotal['Total']) * 100, 2)
age_grouptotal
```

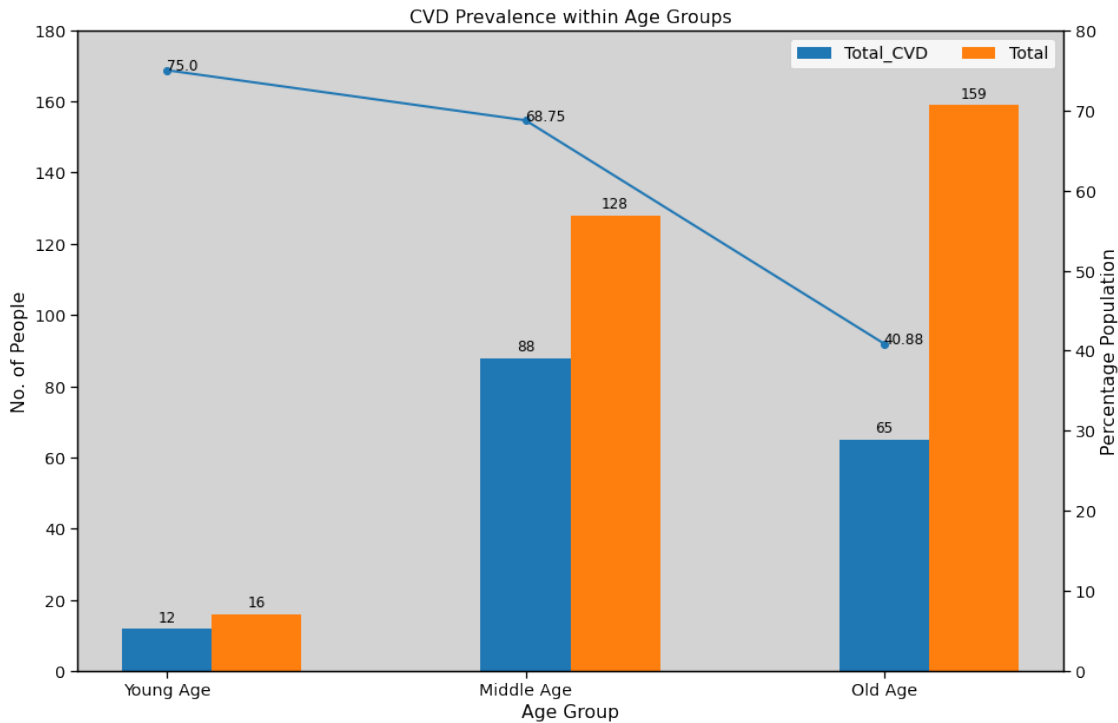
```
[32]:  Age_Category  Total_CVD  Total  Percent_Population
0    Young Age         12      16             75.00
1    Middle Age         88     128             68.75
2     Old Age          65     159             40.88
```

```
[33]: # Create a combined multiple bar plot and line plot to visualize the data of
      ↪ CVD prevalence
```

```

age_cat = ('Young Age', 'Middle Age', 'Old Age') # Label location
cvd_means = {'Total_CVD' : (12, 88, 65),
             'Total' : (16, 128, 159)}           # Define the labels
pp_means = {'Percent_Population' : (75, 68.8, 40.9)}
x = np.arange(len(age_cat))
width = 0.25                                     # Width of the bar
multiplier = 0                                   # Multiplier
fig, ax = plt.subplots(figsize = (15, 10))
ax.set_ylim(0,180)
ax.set_facecolor('lightgray')
ax.grid(False)
ax.set_title('CVD Prevalence within Age Groups')
ax.set_xlabel('Age Group')
ax.set_ylabel('No. of People')
ax.set_xticks(x, age_cat, rotation = 0)
for attribute, measurement in cvd_means.items():
    offset = width * multiplier
    rects = ax.bar(x + offset, measurement, width, label = attribute)
    ax.bar_label(rects, padding=3, fontsize = 12)
    multiplier += 1
ax.legend(loc = 'upper right', ncol = 2)
ax2 = ax.twinx()
ax2.plot(age_grouptotal['Percent_Population'].values, linestyle = '-', marker = 'o',
        linewidth = 2.0)
ax2.set_ylim(0,80)
ax2.grid(False)
ax2.set_ylabel('Percentage Population')
for i, j in age_grouptotal.Percent_Population.items():
    ax2.annotate(str(j), xy = (i,j), fontsize = 12)

```

```
[34]: # Inferences:
# Of the 165 total high risk CVD cases:
# 1. 12 out of 16 Young age are CVD high risk - 75%
# 2. 88 out of 128 Middle age are CVD high risk - 68.75%
# 3. 65 out of 159 Old age are CVD high risk - 40.9%
# It is thus interesting to note the following:
# 1. maximum high risk of CVD is with the young population followed by middle_
→age
# 2. Old age people seem to be at lesser risk of CVD
```

```
[35]: # Refer Line 17
# cp - Chest Pain Type - it has a high positive correlation with the target_
→variable
# Analysis of Chest Pain Type -
# 0 - Asympomatic, 1 - Non-anginal, 2 - Atypical Angina, 3 - Typical Angina
# 4 categorical variables
# Let us understand the distribution of the cp variable

cp_freq = heart.groupby('cp').agg({'cp': 'count'})
cp_freq

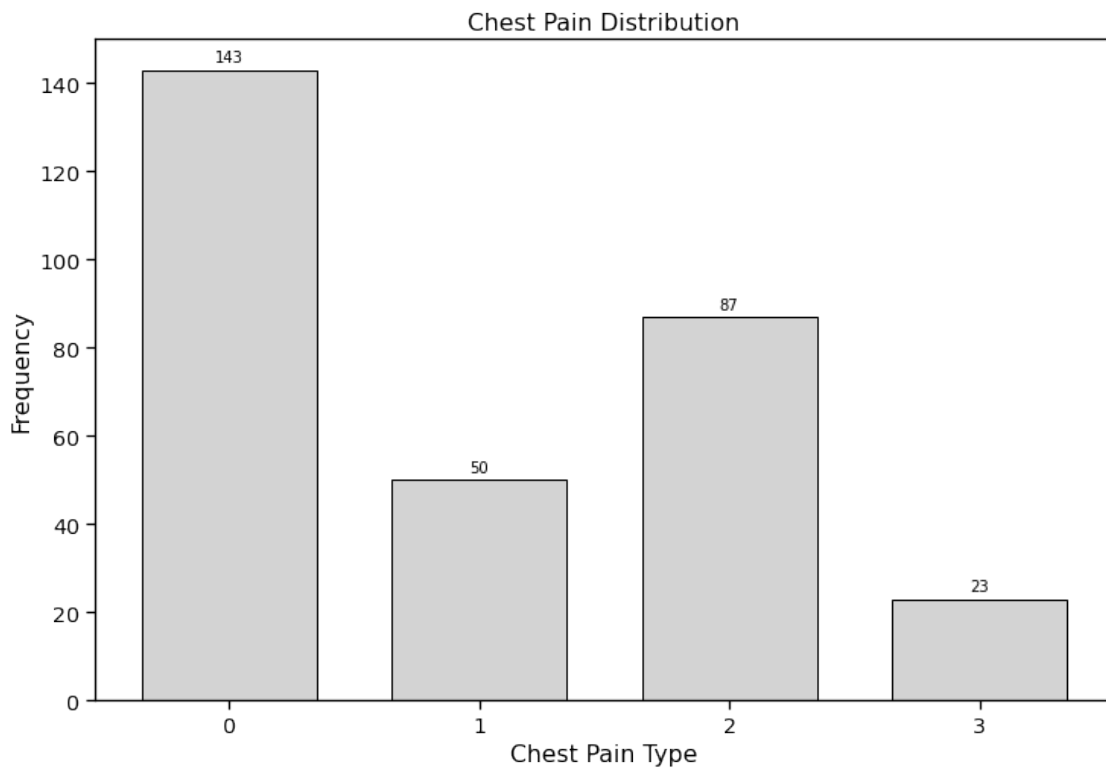
# Inferences
# Of the 303 samples - 143 - asymptomatic, 50 - non-anginal pain, 87 - atypical_
→angina & 23 - typical angina
```

```
[35]:      cp
      cp
0    143
1     50
2     87
3     23
```

```
[36]: # Plot a barplot to see the distribution of samples in the dataset by type of
      ↪ chest pain

plt.figure(figsize = (12,8))
barplot = plt.bar(cp_freq.index, cp_freq.cp, fc = 'lightgray', ec = 'black',
      ↪ width = 0.7)
plt.title('Chest Pain Distribution')
plt.xlabel('Chest Pain Type')
plt.ylabel('Frequency')
plt.ylim([0,150])
plt.xticks(cp_freq.index)
plt.bar_label(barplot, labels = cp_freq.cp, label_type = 'edge', padding = 3,
      ↪ fontsize = 10)
```

```
[36]: [Text(0, 3, '143'), Text(0, 3, '50'), Text(0, 3, '87'), Text(0, 3, '23')]
```



```
[37]: # Understand and check the relation between high risk CVD to the chest pain type

cp_group = heart.groupby('cp')['target'].apply(lambda x: (x == 1).sum()).
    ↪reset_index(name = 'Total_CVD')
cp_group
```

```
[37]:   cp  Total_CVD
0    0         39
1    1         41
2    2         69
3    3         16
```

```
[38]: # Create a new dataframe to display the overall statistics of CVD with type of
    ↪chest pain

cp_grouptotal = pd.DataFrame({'cp' : [0, 1, 2 ,3],
                              'Total_CVD' : [39, 41, 69, 16], 'Total' : [143,
    ↪50, 87, 23]})
cp_grouptotal
```

```
[38]:   cp  Total_CVD  Total
0    0         39   143
1    1         41    50
2    2         69    87
3    3         16    23
```

```
[39]: # Add a % of the population column to the dataframe

cp_grouptotal['Percent_Population'] = round((cp_grouptotal['Total_CVD']/
    ↪cp_grouptotal['Total']) * 100, 2)
cp_grouptotal
```

```
[39]:   cp  Total_CVD  Total  Percent_Population
0    0         39   143             27.27
1    1         41    50             82.00
2    2         69    87             79.31
3    3         16    23             69.57
```

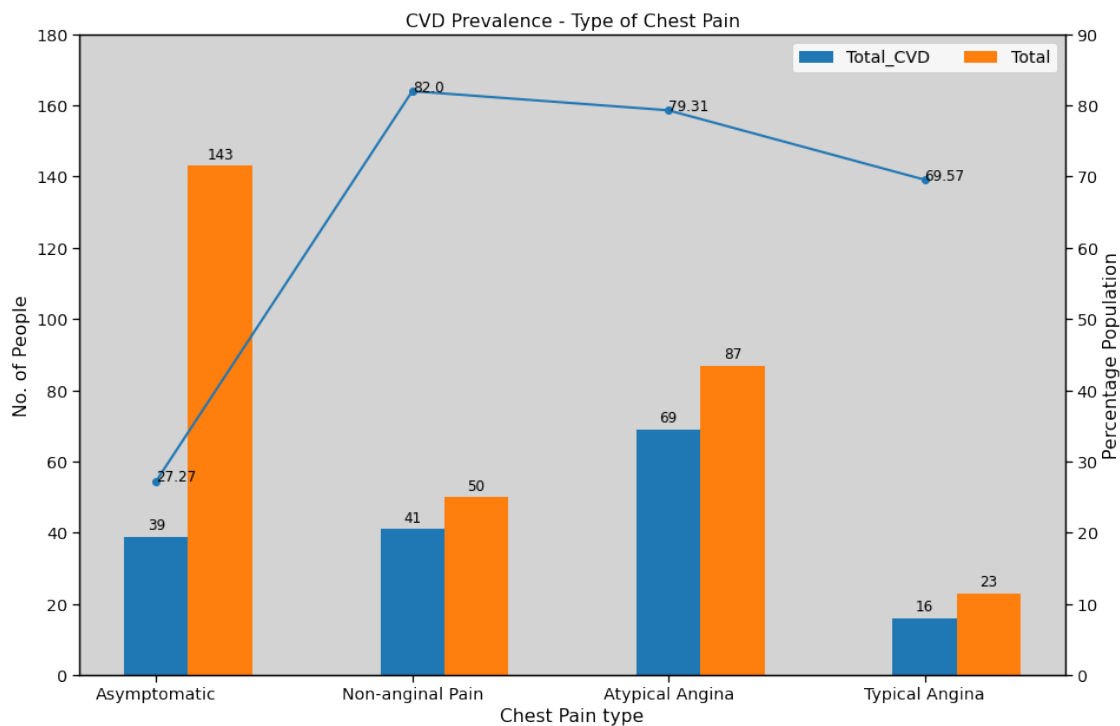
```
[40]: # Create a combined multiple bar plot and line plot to visualize the data of
    ↪CVD prevalence

cp_cat = ('Asymptomatic', 'Non-anginal Pain', 'Atypical Angina', 'Typical
    ↪Angina') # Label location
cvd_means = {'Total_CVD' : (39, 41, 69, 16),
             'Total' : (143, 50, 87, 23)} # Define the labels
pp_means = {'Percent_Population' : (27.3, 82.0, 79.3, 69.6)}
x = np.arange(len(cp_cat))
```

```

width = 0.25                                # Width of the bar
multiplier = 0                               # Multiplier
fig, ax = plt.subplots(figsize = (15, 10))
ax.set_ylim(0,180)
ax.set_facecolor('lightgray')
ax.grid(False)
ax.set_title('CVD Prevalence - Type of Chest Pain')
ax.set_xlabel('Chest Pain type')
ax.set_ylabel('No. of People')
ax.set_xticks(x, cp_cat, rotation = 0)
for attribute, measurement in cvd_means.items():
    offset = width * multiplier
    rects = ax.bar(x + offset, measurement, width, label = attribute)
    ax.bar_label(rects, padding=3, fontsize = 12)
    multiplier += 1
ax.legend(loc = 'upper right', ncol = 2)
ax2 = ax.twinx()
ax2.plot(cp_grouptotal['Percent_Population'].values, linestyle = '--', marker = 'o',
        ↪ 'o', linewidth = 2.0)
ax2.set_ylim(0,90)
ax2.grid(False)
ax2.set_ylabel('Percentage Population')
for i, j in cp_grouptotal.Percent_Population.items():
    ax2.annotate(str(j), xy = (i,j), fontsize = 12)

```



```
[41]: # Inferences:
# Of the 143 total asymptomatic cases, only 39 are high risk CVD cases - 27.3%
# Of the 50 total non-angina pain cases, 41 are high risk CVD cases - 82.0%
# Of the 87 total atypical angina cases, 69 are high risk CVD cases - 79.3%
# Of the 23 total typical angina cases, 23 are high risk CVD cases - 69.8%
# It may be noted that:
# 69.6% to 82.0% of the reported CVD cases are either non-anginal pain,
↳ atypical and typical angina
# Non-anginal pain at 82% has the highest % of positive CVD cases
```

```
[42]: # Refer Line 17
# restecg - resting ecg has a high positive correlation with the target variable
# Analysis of resting ecg - categorical variable with 3 values
# 0 - normal, 1 - having ST-T, 2 - hypertrophy
# Let us understand the distribution of the resting ecg variable

restecg_freq = heart.groupby('restecg').agg({'restecg': 'count'})
restecg_freq

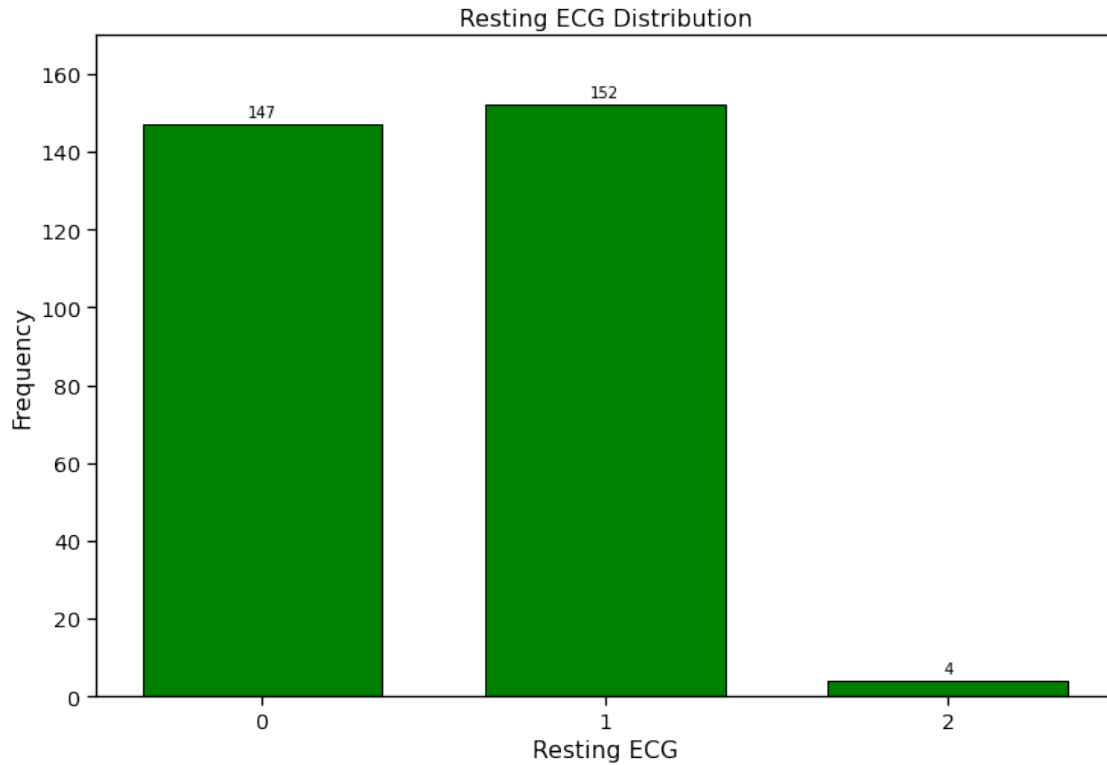
# Inferences
# Of the 303 samples - 147 - normal, 152 - have ST-T and 4 - have hypertrophy
```

```
[42]:          restecg
restecg
0          147
1          152
2           4
```

```
[43]: # Create a barplot to visualize the distribution of the resting ECG data

plt.figure(figsize = (12,8))
barplot = plt.bar(restecg_freq.index, restecg_freq.restecg, fc = 'green', ec =
↳ 'black', width = 0.7)
plt.title('Resting ECG Distribution')
plt.xlabel('Resting ECG')
plt.ylabel('Frequency')
plt.ylim([0,170])
plt.xticks(restecg_freq.index)
plt.bar_label(barplot, labels = restecg_freq.restecg, label_type = 'edge',
↳ padding = 3, fontsize = 10)
```

```
[43]: [Text(0, 3, '147'), Text(0, 3, '152'), Text(0, 3, '4')]
```



```
[44]: # Understand and check the relation between high risk CVD to resting ECG

restecg_group = heart.groupby('restecg')['target'].apply(lambda x: (x == 1).
    ↳sum()).reset_index(name = 'Total_CVD')
restecg_group
```

```
[44]:   restecg  Total_CVD
0         0         68
1         1         96
2         2          1
```

```
[45]: # Create a new dataframe to display the overall statistics of CVD with resting_
    ↳ECG

restecg_grouptotal = pd.DataFrame({'restecg' : [0, 1, 2],
    ↳'Total_CVD' : [68, 96, 1], 'Total' : [147, 152, 4]})
restecg_grouptotal
```

```
[45]:   restecg  Total_CVD  Total
0         0         68    147
1         1         96    152
```

2 2 1 4

```
[46]: # Add a % of the population column to the dataframe

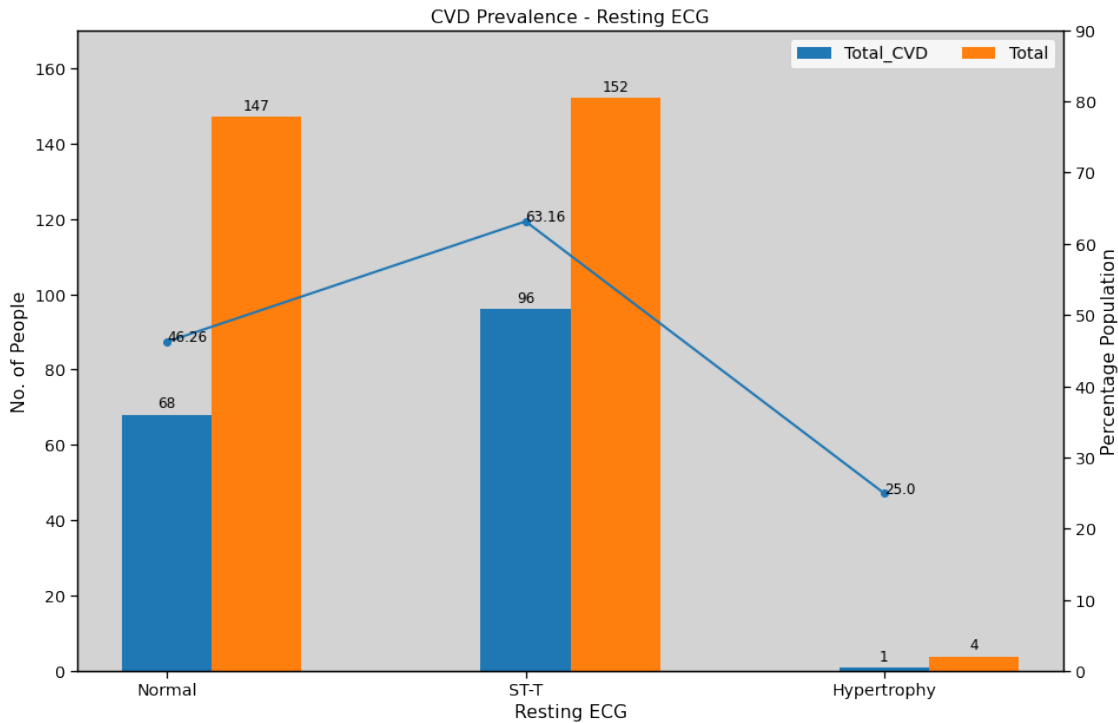
restecg_grouptotal['Percent_Population'] =
    ↪round((restecg_grouptotal['Total_CVD']/
                                restecg_grouptotal['Total']) * 100, 2)
restecg_grouptotal
```

```
[46]:
```

	restecg	Total_CVD	Total	Percent_Population
0	0	68	147	46.26
1	1	96	152	63.16
2	2	1	4	25.00

```
[47]: # Create a combined multiple bar plot and line plot to visualize the data of
    ↪CVD prevalence

restecg_cat = ('Normal', 'ST-T', 'Hypertrophy') # Label location
cvd_means = {'Total_CVD' : (68, 96, 1),
              'Total' : (147, 152, 4)}           # Define the labels
pp_means = {'Percent_Population' : (46.3, 63.2, 25.0)}
x = np.arange(len(restecg_cat))
width = 0.25                                     # Width of the bar
multiplier = 0                                   # Multiplier
fig, ax = plt.subplots(figsize = (15, 10))
ax.set_ylim(0,170)
ax.set_facecolor('lightgray')
ax.grid(False)
ax.set_title('CVD Prevalence - Resting ECG')
ax.set_xlabel('Resting ECG')
ax.set_ylabel('No. of People')
ax.set_xticks(x, restecg_cat, rotation = 0)
for attribute, measurement in cvd_means.items():
    offset = width * multiplier
    rects = ax.bar(x + offset, measurement, width, label = attribute)
    ax.bar_label(rects, padding=3, fontsize = 12)
    multiplier += 1
ax.legend(loc = 'upper right', ncol = 2)
ax2 = ax.twinx()
ax2.plot(restecg_grouptotal['Percent_Population'].values, linestyle = '-',
    ↪marker = 'o', linewidth = 2.0)
ax2.set_ylim(0,90)
ax2.grid(False)
ax2.set_ylabel('Percentage Population')
for i, j in restecg_grouptotal.Percent_Population.items():
    ax2.annotate(str(j), xy = (i,j), fontsize = 12)
```



```
[48]: # Inferences:
# Of the 147 normal ECG cases, 68 are high risk CVD cases - 46.3%
# Of the 152 cases having ST-T, 96 are high risk CVD cases - 63.2%
# Of the 4 cases with Hypertrophy, 1 case is high risk CVD - 25.0%
# It may thus be noted that:
# 1. Maximum high risk of CVD are the ones with ST-T
# 2. The normal resting ecg cases also need to be monitored regularly for
    ↳ inconsistencies
# 3. 46.3% of normal ECG cases also show signs of high risk CVD
```

```
[49]: # Understanding the CVD cases in relation to the sex of the individual

sex_group = heart.groupby('sex')['target'].apply(lambda x: (x == 1).sum()).
    ↳ reset_index(name = 'Total_CVD')
sex_group
```

```
[49]:    sex  Total_CVD
0     0         72
1     1         93
```

```
[50]: # Create a dataframe to analyze the dataset further

sex_grouptotal = pd.DataFrame({'sex' : [0, 1],
                               'Total_CVD' : [72, 93], 'Total' : [96, 207]})
```



```
sex_grouptotal
```

```
[50]:
```

	sex	Total_CVD	Total
0	0	72	96
1	1	93	207

```
[51]: # Add a pecentage of population column to the dataframe
```

```
sex_grouptotal['Percent_Population'] = round((sex_grouptotal['Total_CVD']/  
sex_grouptotal['Total']) * 100, 2)  
sex_grouptotal
```

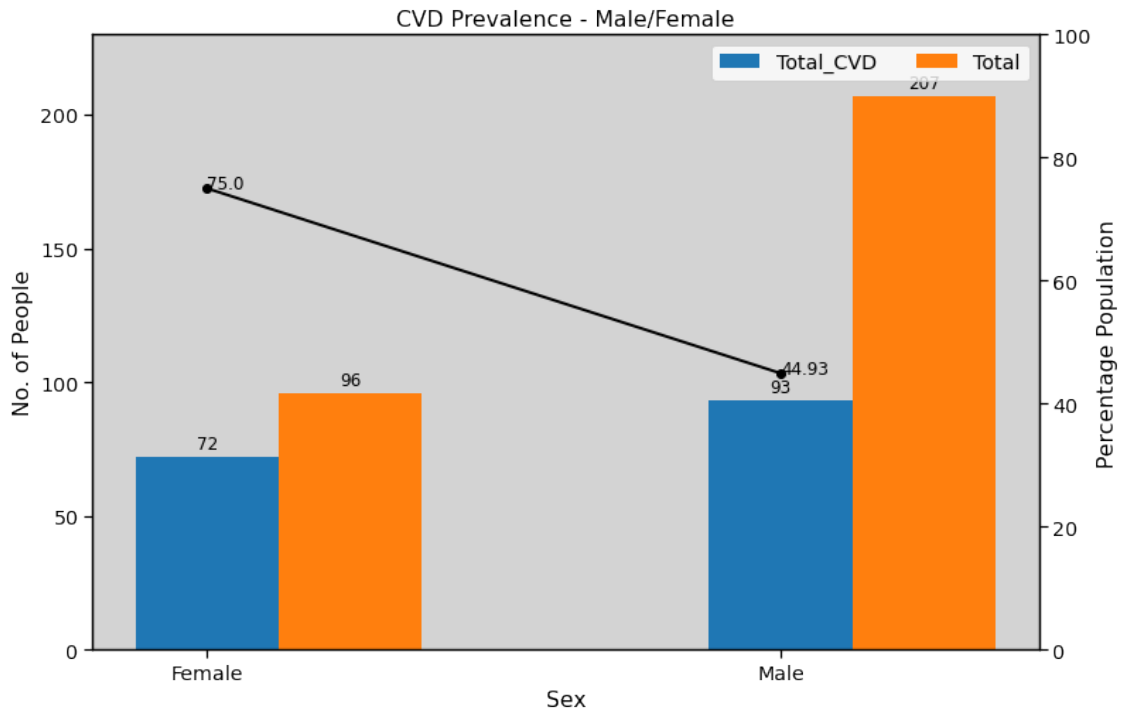
```
[51]:
```

	sex	Total_CVD	Total	Percent_Population
0	0	72	96	75.00
1	1	93	207	44.93

```
[52]: # Create a combined multiple bar plot and line plot to visualize the data of  
↪ CVD prevalence
```

```
sex_cat = ('Female', 'Male') # Label location  
cvd_means = {'Total_CVD' : (72, 93),  
             'Total' : (96, 207)} # Define the labels  
pp_means = {'Percent_Population' : (75.0, 44.9)}  
x = np.arange(len(sex_cat))  
width = 0.25 # Width of the bar  
multiplier = 0 # Multiplier  
fig, ax = plt.subplots(figsize = (12, 8))  
ax.set_ylim(0,230)  
ax.set_facecolor('lightgray')  
ax.grid(False)  
ax.set_title('CVD Prevalence - Male/Female')  
ax.set_xlabel('Sex')  
ax.set_ylabel('No. of People')  
ax.set_xticks(x, sex_cat, rotation = 0)  
for attribute, measurement in cvd_means.items():  
    offset = width * multiplier  
    rects = ax.bar(x + offset, measurement, width, label = attribute)  
    ax.bar_label(rects, padding=3, fontsize = 12)  
    multiplier += 1  
ax.legend(loc = 'upper right', ncol = 2)  
ax2 = ax.twinx()  
ax2.plot(sex_grouptotal['Percent_Population'].values, linestyle = '-', marker = 'o',  
↪ 'o',  
         linewidth = 2.0, color = 'black')  
ax2.set_ylim(0,100)  
ax2.grid(False)  
ax2.set_ylabel('Percentage Population')
```

```
for i, j in sex_grouptotal.Percent_Population.items():
    ax2.annotate(str(j), xy = (i,j), fontsize = 12)
```

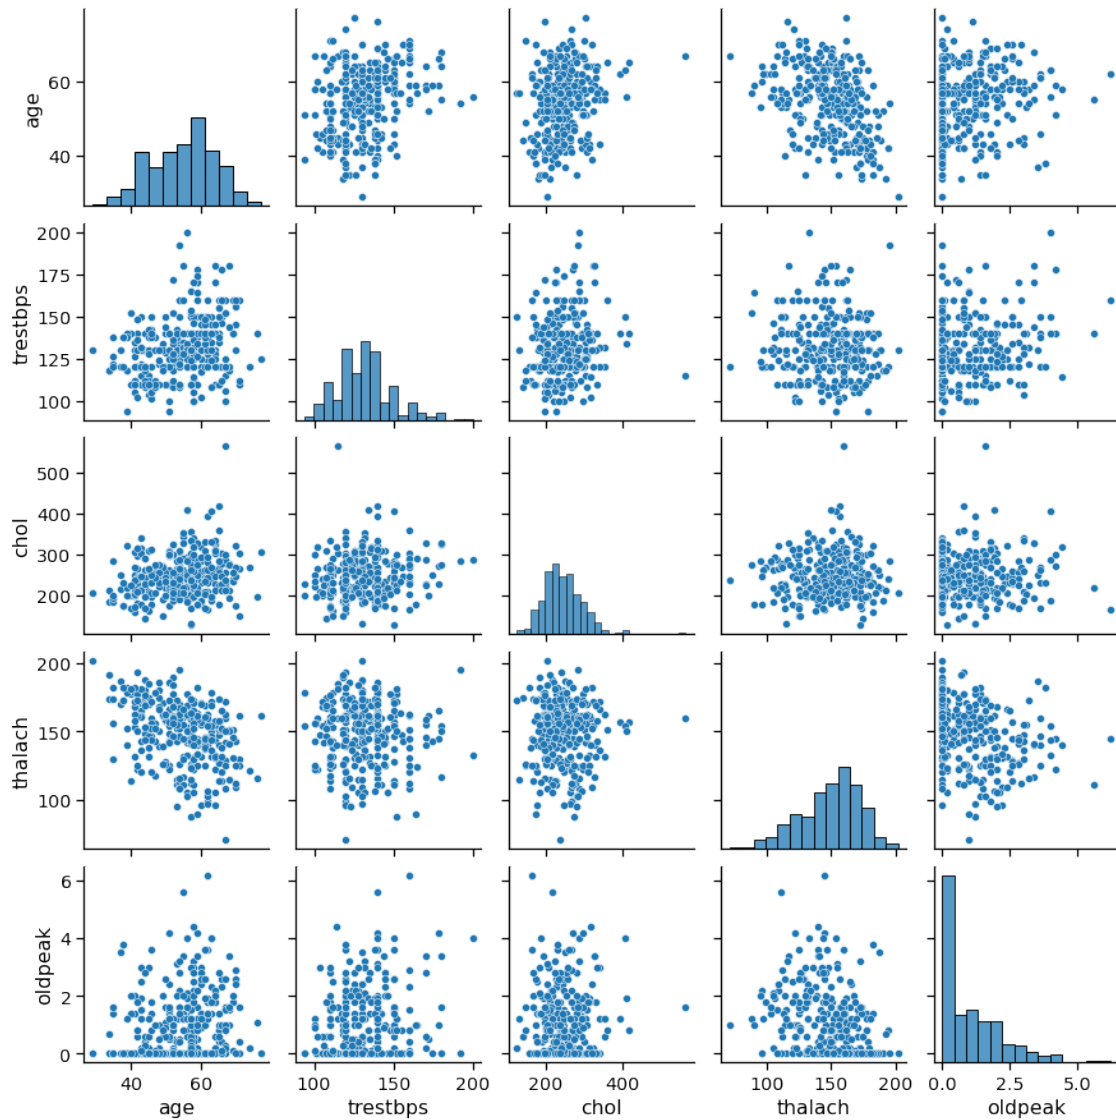


```
[53]: # Inferences:
      # 75% of the Female population are high risk CVD cases
      # 44.6% of the Male population are high risk CVD cases
```

```
[54]: # Create a pair plot to study the correlation of continuous variables

sub_data = heart[['age', 'trestbps', 'chol', 'thalach', 'oldpeak']]
sns.pairplot(sub_data)
```

```
[54]: <seaborn.axisgrid.PairGrid at 0x7f4aac26b2d0>
```



```
[55]: # Refer Line 17
# slope - slope of the peak exercise segment has a high positive correlation
# with the target variable
# Analysis of slope - categorical variable with 3 values
# 0 - upslope, 1 - flat, 2 - downslope
# Let us understand the distribution of the slope variable

slope_freq = heart.groupby('slope').agg({'slope': 'count'})
slope_freq

# Inferences
# Of the 303 samples - 21 - upslope, 140 - have flat and 142 - downslope
```

```
[55]:      slope
slope
0      21
1     140
2     142
```

```
[56]: # Understanding the CVD cases in relation to the ST depression slope

slope_group = heart.groupby('slope')['target'].apply(lambda x: (x == 1).sum()).
↳reset_index(name = 'Total_CVD')
slope_group
```

```
[56]:      slope  Total_CVD
0      0         9
1      1        49
2      2       107
```

```
[57]: # Create a dataframe to analyse the data of ST depression slope further

slope_grouptotal = pd.DataFrame({'slope' : [0, 1, 2],
                                  'Total_CVD' : [9, 49, 107], 'Total' : [21, 140, 142]})
slope_grouptotal
```

```
[57]:      slope  Total_CVD  Total
0      0         9      21
1      1        49     140
2      2       107     142
```

```
[58]: # Add percentage of population column to the dataframe

slope_grouptotal['Percent_Population'] = round((slope_grouptotal['Total_CVD']/
slope_grouptotal['Total']) * 100, 2)
slope_grouptotal
```

```
[58]:      slope  Total_CVD  Total  Percent_Population
0      0         9      21          42.86
1      1        49     140          35.00
2      2       107     142          75.35
```

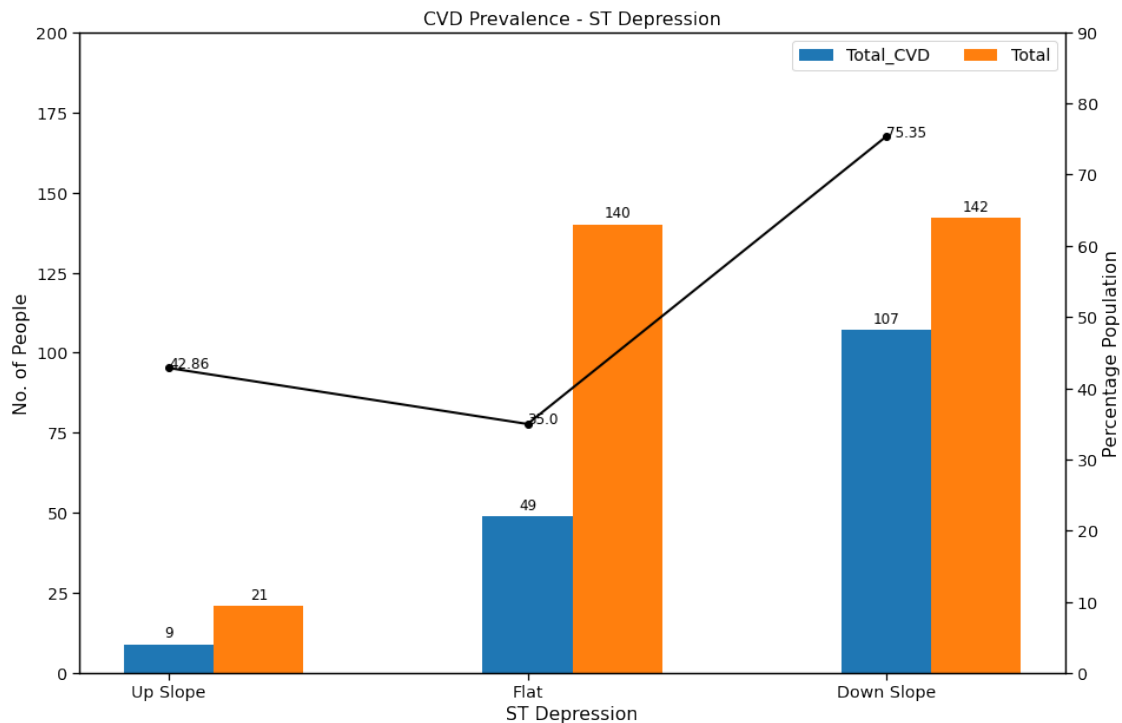
```
[59]: # Create a combined multiple bar plot and line plot to visualize the data of
↳CVD prevalence

slope_cat = ('Up Slope', 'Flat', 'Down Slope') # Label location
cvd_means = {'Total_CVD' : (9, 49, 107),
              'Total' : (21, 140, 142)} # Define the labels
pp_means = {'Percent_Population' : (42.9, 35.0, 75.4)}
```

```

x = np.arange(len(slope_cat))
width = 0.25 # Width of the bar
multiplier = 0 # Multiplier
fig, ax = plt.subplots(figsize = (15, 10))
ax.set_ylim(0,200)
ax.set_facecolor('white')
ax.grid(False)
ax.set_title('CVD Prevalence - ST Depression')
ax.set_xlabel('ST Depression')
ax.set_ylabel('No. of People')
ax.set_xticks(x, slope_cat, rotation = 0)
for attribute, measurement in cvd_means.items():
    offset = width * multiplier
    rects = ax.bar(x + offset, measurement, width, label = attribute)
    ax.bar_label(rects, padding=3, fontsize = 12)
    multiplier += 1
ax.legend(loc = 'upper right', ncol = 2)
ax2 = ax.twinx()
ax2.plot(slope_grouptotal['Percent_Population'].values, linestyle = '-', marker_
    => 'o',
        linewidth = 2.0, color = 'black')
ax2.set_ylim(0,90)
ax2.grid(False)
ax2.set_ylabel('Percentage Population')
for i, j in slope_grouptotal.Percent_Population.items():
    ax2.annotate(str(j), xy = (i,j), fontsize = 12)

```




```
[60]: # Inferences
# slope - CVD instances are highest for ST depression down slope cases at 75.4%
# of the cases
# slope - CVD instances are below 50% for ST depression up slope and and flat
```

```
[61]: # Refer Line 17
# thalach - maximum heart rate achieved has a high positive correlation with
# the target variable
# Analysis of thalach variable
# Let us understand the distribution of the maximum heart rate achieved variable

sns.displot(heart.thalach, color = 'red', kde = True, height = 6, aspect = 2)
plt.title('Maximum Heart Rate Achieved')
plt.xlabel('Heart Rate')
plt.ylabel('Density')
plt.ylim(0, 70)
```

```
[61]: (0.0, 70.0)
```




```
[62]: # Understand the distribution of the data of highest heart rate achieved
# Study the data in relation to the target variable and sex of the individual
# Use a box plot to understand the min, max, median and outliers in the dataset

plt.figure(figsize = (12, 8))
sns.boxplot(x = 'target', y = 'thalach', hue = 'sex', data = heart)
plt.title('Maximum Heart Rate v/s Heart Disease', fontsize = 20)
```

```
plt.xlabel('Heart Disease Risk', fontsize = 16)
plt.ylabel('Maximum Heart Rate Reached', fontsize = 16)
```

```
[62]: Text(0, 0.5, 'Maximum Heart Rate Reached')
```



```
[63]: # Inferences
# High risk CVD (positive patients) shows a higher median
# The median values for male and female is not very different in both cases
# Males have a higher range of maximum heart rate reached than females
```

```
[64]: # Refer Line 17
# There are 2 variables which have a high negative correlation to the target_
↪ variable
# exang - exercise induced angina - (-) 0.44
# oldpeak - ST depression induced by exercise - (-) 0.43
# We need to analyze these 2 variables in more detail
# 1. exang - 1 - yes, 0 - no

exang_freq = heart.groupby('exang').agg({'exang': 'count'})
exang_freq

# Inferences
```

```
# Of the 303 samples - 204 - No exercise induced angina, 99 - Yes to exercise_  
→induced angina
```

```
[64]:      exang  
exang  
0      204  
1       99
```

```
[65]: # Understanding the CVD cases in relation to the exercise induced angina  
  
exang_group = heart.groupby('exang')['target'].apply(lambda x: (x == 1).sum()).  
→reset_index(name = 'Total_CVD')  
exang_group
```

```
[65]:      exang  Total_CVD  
0      0      142  
1      1       23
```

```
[66]: # Create a dataframe to analyse the data of exercise induced angina further  
  
exang_grouptotal = pd.DataFrame({'exang' : [0, 1],  
                                'Total_CVD' : [142, 23], 'Total' : [204, 99]})  
exang_grouptotal
```

```
[66]:      exang  Total_CVD  Total  
0      0      142      204  
1      1       23       99
```

```
[67]: # Add percentage of population column to the dataframe  
  
exang_grouptotal['Percent_Population'] = round((exang_grouptotal['Total_CVD']/  
                                                exang_grouptotal['Total']) * 100, 2)  
exang_grouptotal
```

```
[67]:      exang  Total_CVD  Total  Percent_Population  
0      0      142      204          69.61  
1      1       23       99          23.23
```

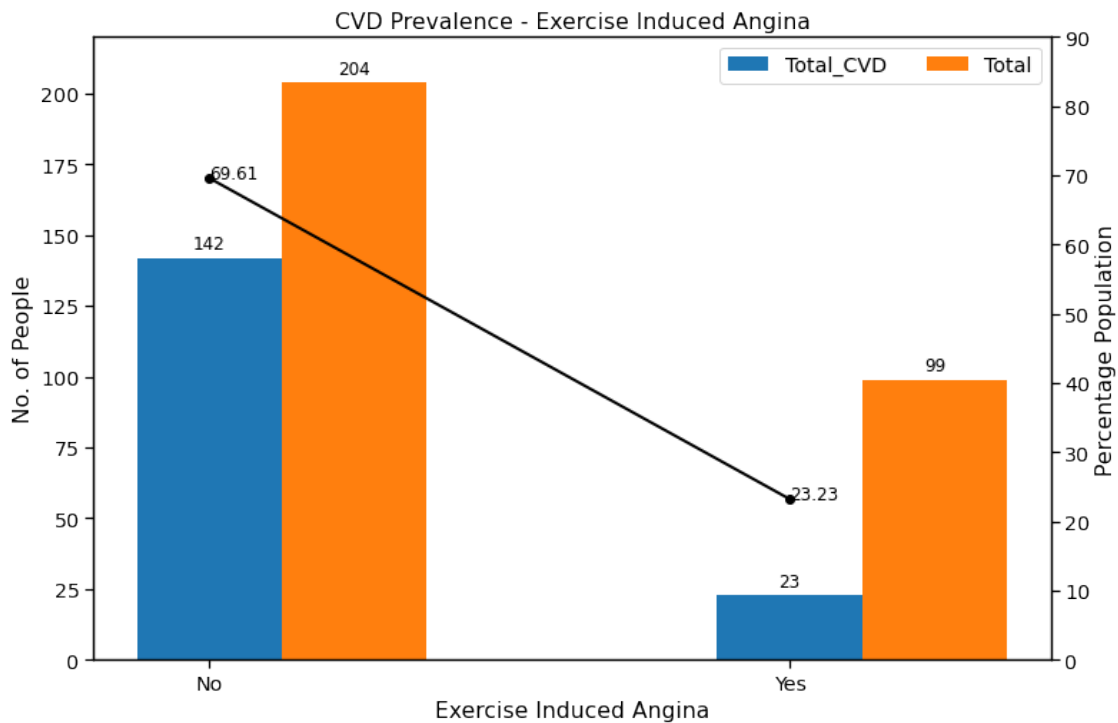
```
[68]: # Create a combined multiple bar plot and line plot to visualize the data of_  
→CVD prevalence  
  
exang_cat = ('No', 'Yes') # Label location  
cvd_means = {'Total_CVD' : (142, 23),  
             'Total' : (204, 99)} # Define the labels  
pp_means = {'Percent_Population' : (42.9, 35.0, 75.4)}  
x = np.arange(len(exang_cat))  
width = 0.25 # Width of the bar
```



```

multiplier = 0                                # Multiplier
fig, ax = plt.subplots(figsize = (12, 8))
ax.set_ylim(0,220)
ax.set_facecolor('white')
ax.grid(False)
ax.set_title('CVD Prevalence - Exercise Induced Angina')
ax.set_xlabel('Exercise Induced Angina')
ax.set_ylabel('No. of People')
ax.set_xticks(x, exang_cat, rotation = 0)
for attribute, measurement in cvd_means.items():
    offset = width * multiplier
    rects = ax.bar(x + offset, measurement, width, label = attribute)
    ax.bar_label(rects, padding=3, fontsize = 12)
    multiplier += 1
ax.legend(loc = 'upper right', ncol = 2)
ax2 = ax.twinx()
ax2.plot(exang_grouptotal['Percent_Population'].values, linestyle = '-', marker_
    ↳ = 'o',
        linewidth = 2.0, color = 'black')
ax2.set_ylim(0,90)
ax2.grid(False)
ax2.set_ylabel('Percentage Population')
for i, j in exang_grouptotal.Percent_Population.items():
    ax2.annotate(str(j), xy = (i,j), fontsize = 12)

```

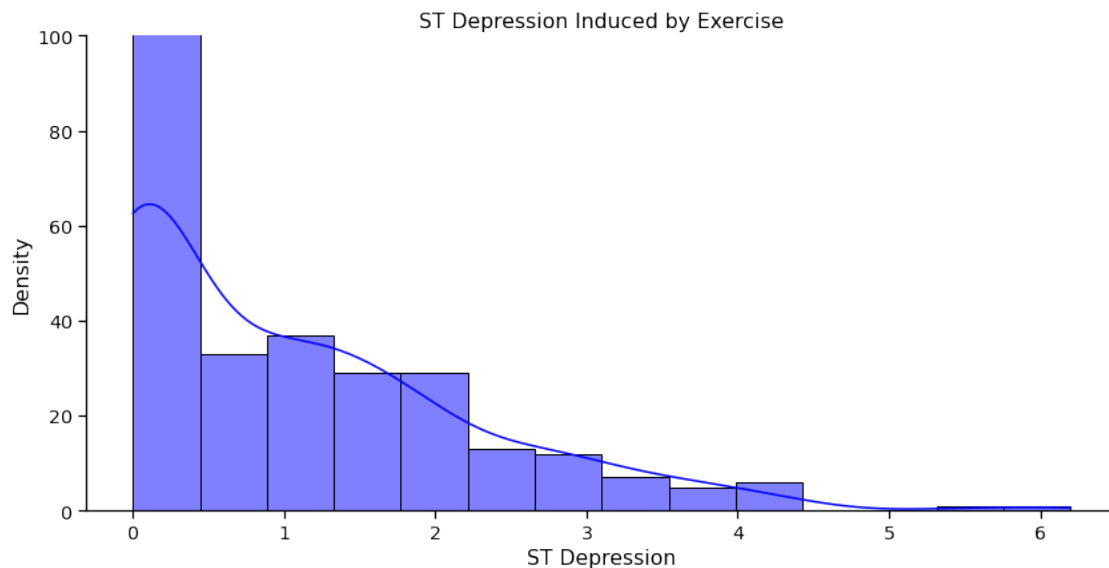


```
[69]: # Inferences
# Of the total positive cases, there are only 23.2% of cases where there has
↳ been exercise induced angina
```

```
[70]: # Refer Line 64
# 2. oldpeak - ST depression induced by exercise relative to rest - (-) 0.43
# Continuous variable
# Let us understand the distribution of the maximum oldpeak variable

sns.displot(heart.oldpeak, color = 'blue', kde = True, height = 6, aspect = 2)
plt.title('ST Depression Induced by Exercise')
plt.xlabel('ST Depression')
plt.ylabel('Density')
plt.ylim(0, 100)
```

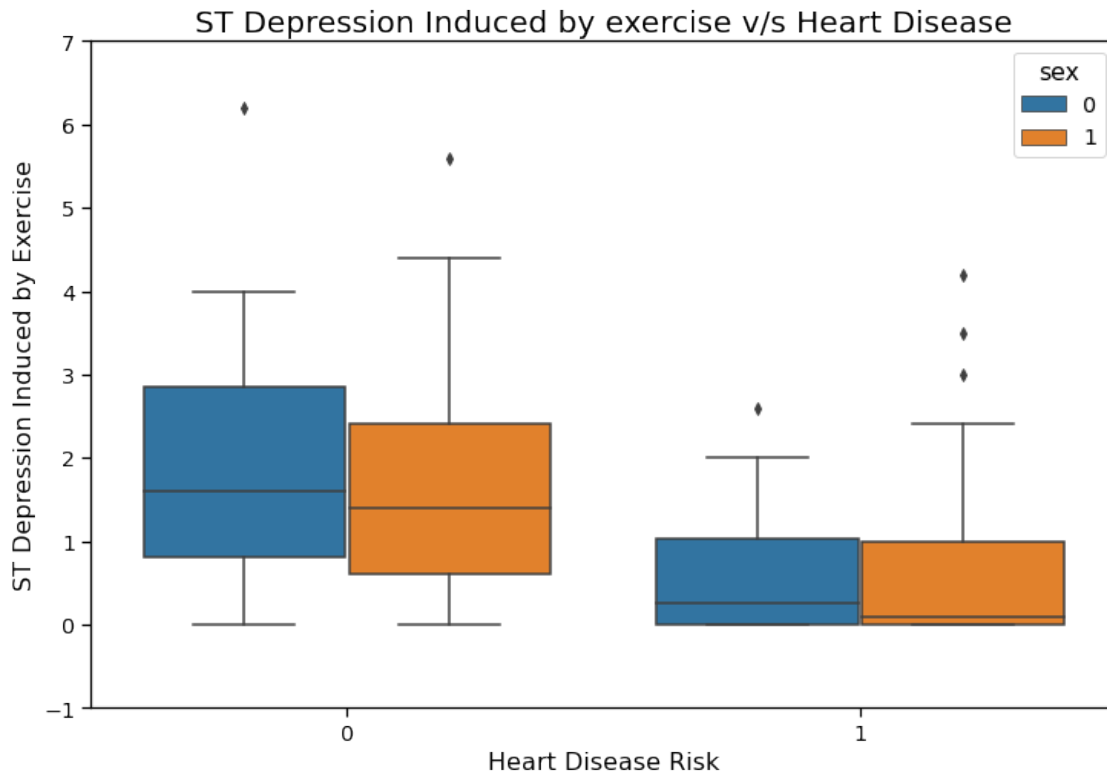
```
[70]: (0.0, 100.0)
```



```
[71]: # Understand the distribution of the data of ST depression induced by exercise
# Study the data in relation to the target variable and sex of the individual
# Use a box plot to understand the min, max, median and outliers in the dataset

plt.figure(figsize = (12, 8))
sns.boxplot(x = 'target', y = 'oldpeak', hue = 'sex', data = heart)
plt.title('ST Depression Induced by exercise v/s Heart Disease', fontsize = 20)
plt.xlabel('Heart Disease Risk', fontsize = 16)
plt.ylabel('ST Depression Induced by Exercise', fontsize = 16)
plt.ylim(-1, 7)
```

[71]: (-1.0, 7.0)



```
[72]: # Inferences
# In positive CVD cases, the median of ST depression induced by exercise is
# much low
# Lower ST depression induced by exercise median is indicative of heart disease
# In negative CVD cases, the median of ST depression induced by exercise for
# males and females is almost same
```

```
[73]: # filter CVD positive cases for further analysis
```

```
pos_data = heart[heart['target'] == 1]
pos_data.describe()
```

```
[73]:
```

	age	sex	cp	trestbps	chol	fbs \
count	165.000000	165.000000	165.000000	165.000000	165.000000	165.000000
mean	52.496970	0.563636	1.375758	129.303030	242.230303	0.139394
std	9.550651	0.497444	0.952222	16.169613	53.552872	0.347412
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000
25%	44.000000	0.000000	1.000000	120.000000	208.000000	0.000000
50%	52.000000	1.000000	2.000000	130.000000	234.000000	0.000000
75%	59.000000	1.000000	2.000000	140.000000	267.000000	0.000000

max	76.000000	1.000000	3.000000	180.000000	564.000000	1.000000
-----	-----------	----------	----------	------------	------------	----------

	restecg	thalach	exang	oldpeak	slope	ca \
count	165.000000	165.000000	165.000000	165.000000	165.000000	165.000000
mean	0.593939	158.466667	0.139394	0.583030	1.593939	0.363636
std	0.504818	19.174276	0.347412	0.780683	0.593635	0.848894
min	0.000000	96.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	149.000000	0.000000	0.000000	1.000000	0.000000
50%	1.000000	161.000000	0.000000	0.200000	2.000000	0.000000
75%	1.000000	172.000000	0.000000	1.000000	2.000000	0.000000
max	2.000000	202.000000	1.000000	4.200000	2.000000	4.000000

	thal	target
count	165.000000	165.0
mean	2.121212	1.0
std	0.465752	0.0
min	0.000000	1.0
25%	2.000000	1.0
50%	2.000000	1.0
75%	2.000000	1.0
max	3.000000	1.0

```
[74]: # filter CVD positive cases for further analysis
```

```
neg_data = heart[heart['target'] == 0]
neg_data.describe()
```

```
[74]:
```

	age	sex	cp	trestbps	chol	fbs \
count	138.000000	138.000000	138.000000	138.000000	138.000000	138.000000
mean	56.601449	0.826087	0.478261	134.398551	251.086957	0.159420
std	7.962082	0.380416	0.905920	18.729944	49.454614	0.367401
min	35.000000	0.000000	0.000000	100.000000	131.000000	0.000000
25%	52.000000	1.000000	0.000000	120.000000	217.250000	0.000000
50%	58.000000	1.000000	0.000000	130.000000	249.000000	0.000000
75%	62.000000	1.000000	0.000000	144.750000	283.000000	0.000000
max	77.000000	1.000000	3.000000	200.000000	409.000000	1.000000

	restecg	thalach	exang	oldpeak	slope	ca \
count	138.000000	138.000000	138.000000	138.000000	138.000000	138.000000
mean	0.449275	139.101449	0.550725	1.585507	1.166667	1.166667
std	0.541321	22.598782	0.499232	1.300340	0.561324	1.043460
min	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	125.000000	0.000000	0.600000	1.000000	0.000000
50%	0.000000	142.000000	1.000000	1.400000	1.000000	1.000000
75%	1.000000	156.000000	1.000000	2.500000	1.750000	2.000000
max	2.000000	195.000000	1.000000	6.200000	2.000000	4.000000

	thal	target
count	138.000000	138.0
mean	2.543478	0.0
std	0.684762	0.0
min	0.000000	0.0
25%	2.000000	0.0
50%	3.000000	0.0
75%	3.000000	0.0
max	3.000000	0.0

```
[75]: heart.columns
```

```
[75]: Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach',
        'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target', 'Age_Category'],
        dtype='object')
```

```
[76]: # Creating a dataframe to compare the mean of the positive CVD and negative CVD
        ↪ filtered data above

compare_data = pd.DataFrame({'Health Variable' : ['trestbps', 'chol',
        ↪ 'thalach', 'oldpeak'],
                             'Pos_CVD_Mean' : [129.3, 242.2, 158.5, 0.58],
                             'Neg_CVD_Mean' : [134.4, 251.1, 139.1, 1.58]})

compare_data
```

```
[76]:   Health Variable  Pos_CVD_Mean  Neg_CVD_Mean
0      trestbps         129.30         134.40
1         chol         242.20         251.10
2      thalach         158.50         139.10
3      oldpeak           0.58           1.58
```

```
[77]: # Inferences
        # In positive cases, the highest heart rate achieved is much higher than
        ↪ negative cases
        # In positive cases, the ST depression induced by exercise is very low compared
        ↪ to negative cases
        # Ratio is almost 1/3rd for positive cases in terms of ST depression induced by
        ↪ exercise
```

```
[78]: heart1 = heart.drop(['Age_Category'], axis = 1)
        heart1
```

```
[78]:   age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  \
0    63    1   3     145    233    1         0     150      0      2.3
1    37    1   2     130    250    0         1     187      0      3.5
2    41    0   1     130    204    0         0     172      0      1.4
3    56    1   1     120    236    0         1     178      0      0.8
```

4	57	0	0	120	354	0	1	163	1	0.6
..
298	57	0	0	140	241	0	1	123	1	0.2
299	45	1	3	110	264	0	1	132	0	1.2
300	68	1	0	144	193	1	1	141	0	3.4
301	57	1	0	130	131	0	1	115	1	1.2
302	57	0	1	130	236	0	0	174	0	0.0

	slope	ca	thal	target
0	0	0	1	1
1	0	0	2	1
2	2	0	2	1
3	2	0	2	1
4	2	0	2	1
..
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

[303 rows x 14 columns]

```
[79]: # Model building - Logistic Regression
# Import packages

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
```

```
[80]: # Assigning the 13 variables to X and the last column to Predictor Y

x = heart.drop(['target', 'Age_Category'], axis = 1)
x
```

```
[80]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63	1	3	145	233	1	0	150	0	2.3	
1	37	1	2	130	250	0	1	187	0	3.5	
2	41	0	1	130	204	0	0	172	0	1.4	
3	56	1	1	120	236	0	1	178	0	0.8	
4	57	0	0	120	354	0	1	163	1	0.6	
..
298	57	0	0	140	241	0	1	123	1	0.2	
299	45	1	3	110	264	0	1	132	0	1.2	
300	68	1	0	144	193	1	1	141	0	3.4	
301	57	1	0	130	131	0	1	115	1	1.2	
302	57	0	1	130	236	0	0	174	0	0.0	

	slope	ca	thal
0	0	0	1
1	0	0	2
2	2	0	2
3	2	0	2
4	2	0	2
..
298	1	0	3
299	1	0	3
300	1	2	3
301	1	1	3
302	1	1	2

[303 rows x 13 columns]

```
[81]: y = heart.target
      y
```

```
[81]: 0      1
      1      1
      2      1
      3      1
      4      1
      ..
      298    0
      299    0
      300    0
      301    0
      302    0
      Name: target, Length: 303, dtype: int64
```

```
[82]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.2,
      ↪random_state = 1)
```

```
[83]: df_train = pd.concat([x_train, y_train], axis = 1)
      df_train
```

```
[83]:   age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  \
62   52   1   3     118    186   0         0     190     0     0.0
127  67   0   2     152    277   0         1     172     0     0.0
111  57   1   2     150    126   1         1     173     0     0.2
287  57   1   1     154    232   0         0     164     0     0.0
108  50   0   1     120    244   0         1     162     0     1.1
..   ..   ..   ..   ..   ..   ..   ..   ..   ..   ..
203  68   1   2     180    274   1         0     150     1     1.6
255  45   1   0     142    309   0         0     147     1     0.0
```

72	29	1	1	130	204	0	0	202	0	0.0
235	51	1	0	140	299	0	1	173	1	1.6
37	54	1	2	150	232	0	0	165	0	1.6

	slope	ca	thal	target
62	1	0	1	1
127	2	1	2	1
111	2	1	3	1
287	2	1	2	0
108	2	0	2	1
..
203	1	0	3	0
255	1	3	3	0
72	2	0	2	1
235	2	0	3	0
37	2	0	3	1

[242 rows x 14 columns]

```
[84]: df_test = pd.concat([x_test, y_test], axis = 1)
df_test
```

```
[84]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
204	62	0	0	160	164	0	0	145	0	6.2	
159	56	1	1	130	221	0	0	163	0	0.0	
219	48	1	0	130	256	1	0	150	1	0.0	
174	60	1	0	130	206	0	0	132	1	2.4	
184	50	1	0	150	243	0	0	128	0	2.6	
..			
0	63	1	3	145	233	1	0	150	0	2.3	
288	57	1	0	110	335	0	1	143	1	3.0	
259	38	1	3	120	231	0	1	182	1	3.8	
179	57	1	0	150	276	0	0	112	1	0.6	
110	64	0	0	180	325	0	1	154	1	0.0	

	slope	ca	thal	target
204	0	3	3	0
159	2	0	3	1
219	2	2	3	0
174	1	2	3	0
184	1	0	3	0
..
0	0	0	1	1
288	1	1	3	0
259	1	0	3	0
179	1	1	1	0
110	2	0	2	1

[61 rows x 14 columns]

```
[85]: x_train.shape
```

```
[85]: (242, 13)
```

```
[86]: y_train.shape
```

```
[86]: (242,)
```

```
[87]: x_test.shape
```

```
[87]: (61, 13)
```

```
[88]: y_test.shape
```

```
[88]: (61,)
```

```
[89]: # Normalize the train and test data prior to regression  
# Import the package StandardScaler from sklearn  
  
from sklearn.preprocessing import StandardScaler
```

```
[90]: # Normalize the x_train and x_test data  
  
sc = StandardScaler()  
x_train = sc.fit_transform(x_train)  
x_test = sc.transform(x_test)
```

```
[91]: # Logistic Regression  
# Get instance of the model & fit the train model  
  
model1 = LogisticRegression(random_state = 1)  
model1.fit(x_train, y_train)
```

```
[91]: LogisticRegression(random_state=1)
```

```
[92]: # Get the y predictions  
# print the model accuracy classification report  
  
y_pred1 = model1.predict(x_test)  
print(classification_report(y_test, y_pred1))
```

	precision	recall	f1-score	support
0	0.77	0.67	0.71	30

1	0.71	0.81	0.76	31
accuracy			0.74	61
macro avg	0.74	0.74	0.74	61
weighted avg	0.74	0.74	0.74	61

```
[93]: # Logistic Regression accuracy is 74%
```

```
[94]: # K-NN - K-Nearest Neighbors
# Import package

from sklearn.neighbors import KNeighborsClassifier
```

```
[95]: # Get instance of the model & fit the train model

model2 = KNeighborsClassifier()
model2.fit(x_train, y_train)
```

```
[95]: KNeighborsClassifier()
```

```
[96]: # Get the y predictions
# print the model accuracy classification report

y_pred2 = model2.predict(x_test)
print(classification_report(y_test, y_pred2))
```

	precision	recall	f1-score	support
0	0.78	0.70	0.74	30
1	0.74	0.81	0.77	31
accuracy			0.75	61
macro avg	0.76	0.75	0.75	61
weighted avg	0.76	0.75	0.75	61

```
[97]: # K-NN - K-Nearest Neighbors accuracy - 75%
```

```
[98]: # Support Vector Machine
# Import package

from sklearn.svm import SVC
```

```
[99]: # Get instance of the model & fit the train model

model3 = SVC(random_state = 1)
```

```
model3.fit(x_train, y_train)
```

```
[99]: SVC(random_state=1)
```

```
[100]: # Get the y predictions  
# print the model accuracy classification report  
  
y_pred3 = model3.predict(x_test)  
print(classification_report(y_test, y_pred3))
```

	precision	recall	f1-score	support
0	0.80	0.67	0.73	30
1	0.72	0.84	0.78	31
accuracy			0.75	61
macro avg	0.76	0.75	0.75	61
weighted avg	0.76	0.75	0.75	61

```
[101]: # Support Vector Machine SVC accuracy is 75%
```

```
[102]: # Naives Bayer Classifier  
# Import package  
  
from sklearn.naive_bayes import GaussianNB
```

```
[103]: # Get instance of the model & fit the train model  
  
model4 = GaussianNB()  
model4.fit(x_train, y_train)
```

```
[103]: GaussianNB()
```

```
[104]: # Get the y predictions  
# print the model accuracy classification report  
  
y_pred4 = model4.predict(x_test)  
print(classification_report(y_test, y_pred4))
```

	precision	recall	f1-score	support
0	0.79	0.73	0.76	30
1	0.76	0.81	0.78	31
accuracy			0.77	61
macro avg	0.77	0.77	0.77	61

weighted avg	0.77	0.77	0.77	61
--------------	------	------	------	----

```
[105]: # Naives Bayer Classifier accuracy - 77%
```

```
[106]: # Decision Tree
# Import package

from sklearn.tree import DecisionTreeClassifier
```

```
[107]: # Get instance of the model & fit the train model

model5 = DecisionTreeClassifier(random_state = 1)
model5.fit(x_train, y_train)
```

```
[107]: DecisionTreeClassifier(random_state=1)
```

```
[108]: # Get the y predictions
# print the model accuracy classification report

y_pred5 = model5.predict(x_test)
print(classification_report(y_test, y_pred5))
```

	precision	recall	f1-score	support
0	0.68	0.70	0.69	30
1	0.70	0.68	0.69	31
accuracy			0.69	61
macro avg	0.69	0.69	0.69	61
weighted avg	0.69	0.69	0.69	61

```
[109]: # Decision Tree Classifier accuracy is 69%
```

```
[110]: # Random Forest
# Import package

from sklearn.ensemble import RandomForestClassifier
```

```
[111]: # Get instance of the model & fit the train model

model6 = RandomForestClassifier(random_state = 1)
model6.fit(x_train, y_train)
```

```
[111]: RandomForestClassifier(random_state=1)
```

```
[112]: # Get the y predictions
# print the model accuracy classification report
```

```
y_pred6 = model6.predict(x_test)
print(classification_report(y_test, y_pred6))
```

	precision	recall	f1-score	support
0	0.88	0.70	0.78	30
1	0.76	0.90	0.82	31
accuracy			0.80	61
macro avg	0.82	0.80	0.80	61
weighted avg	0.81	0.80	0.80	61

```
[113]: # Random Forest Classifier accuracy is 80%
```

```
[114]: # Random Forest Classifier accuracy is the highest of all models - 80%
```

```
[115]: # Confusion Matrix
# Import package

from sklearn.metrics import confusion_matrix, accuracy_score
```

```
[116]: cm = confusion_matrix(y_test, y_pred6)
cm
```

```
[116]: array([[21,  9],
        [ 3, 28]])
```

```
[117]: accuracy_score(y_test, y_pred6)
```

```
[117]: 0.8032786885245902
```

```
[118]: # Accuracy of 80.32% is good as any accuracy > 70% is good to deploy a model
# Inferences from Confusion Matrix
# True Positives (TP) - 21
# True Negatives (TN) - 28
# No. of errors - 9 & 3
# Type 1 error - False Positives (FP) - 9 - Predicted positive and its false
# Type 2 error - False Negatives (FN) - 3 - Predicted negative and its false
# Calculate accuracy - (TP + TN) / (TP + TN + FP + FN) - Correct Predicted /
↳ Total
# Accuracy = (21 + 28) / (21 + 28 + 9 + 3)
# Accuracy - 49 / 61 = 80.32%
```

```
[119]: # Feature Importance
```

```
importance = model6.feature_importances_
```

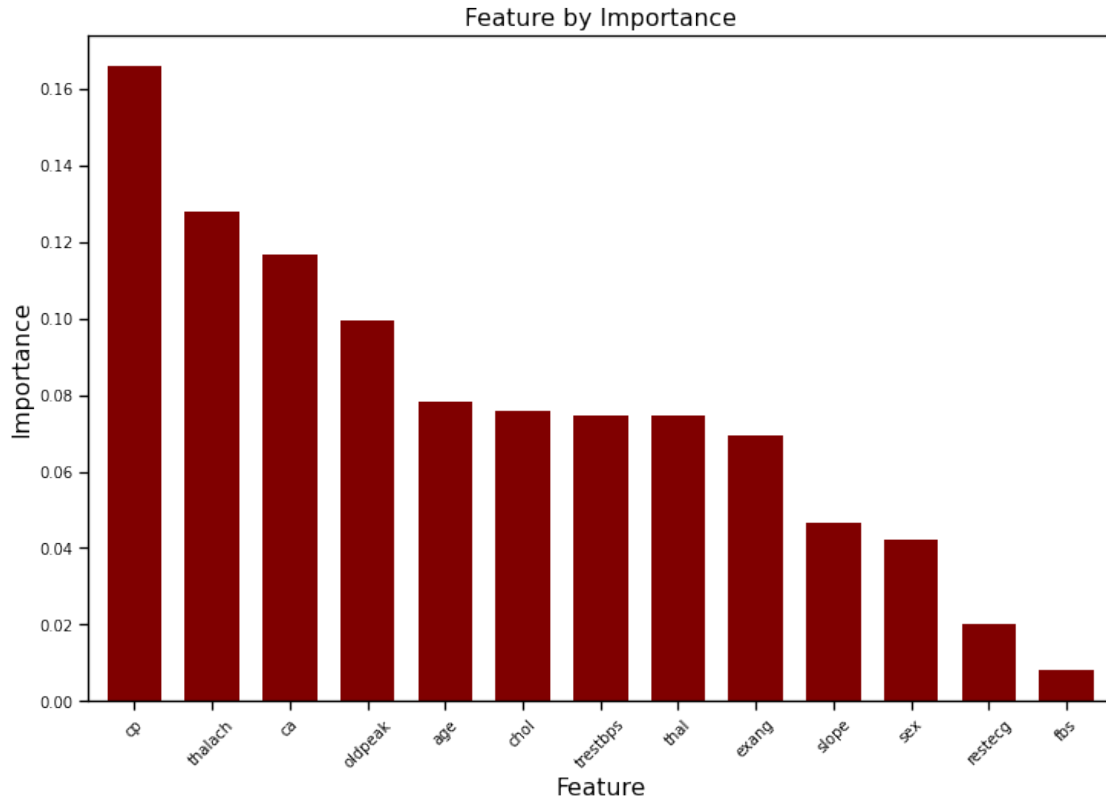
```
[120]: # Summarize feature importance
```

```
for i,v in enumerate(importance):  
    print('Feature: %0d, Score: %.5f' % (i,v))
```

```
Feature: 0, Score: 0.07814  
Feature: 1, Score: 0.04206  
Feature: 2, Score: 0.16580  
Feature: 3, Score: 0.07477  
Feature: 4, Score: 0.07587  
Feature: 5, Score: 0.00828  
Feature: 6, Score: 0.02014  
Feature: 7, Score: 0.12772  
Feature: 8, Score: 0.06950  
Feature: 9, Score: 0.09957  
Feature: 10, Score: 0.04677  
Feature: 11, Score: 0.11667  
Feature: 12, Score: 0.07473
```

```
[121]: index = x.columns  
importance = pd.Series(model6.feature_importances_, index = index)  
importance.nlargest(13).plot(kind = 'bar', color = 'maroon', width = 0.7,  
    ↳figsize = (12, 8))  
plt.title('Feature by Importance')  
plt.xlabel('Feature')  
plt.ylabel('Importance')  
plt.xticks(fontsize = 10, rotation = 45)  
plt.yticks(fontsize = 10)
```

```
[121]: (array([0. , 0.02, 0.04, 0.06, 0.08, 0.1 , 0.12, 0.14, 0.16, 0.18]),  
    [Text(0, 0, ''),  
      Text(0, 0, ''),  
      Text(0, 0, ''),  
      Text(0, 0, ''),  
      Text(0, 0, ''),  
      Text(0, 0, ''),  
      Text(0, 0, ''),  
      Text(0, 0, ''),  
      Text(0, 0, ''),  
      Text(0, 0, ''),  
      Text(0, 0, '')[12]])
```



```
[122]: # Inferences
# The top 4 significant features that determine/predict CVD positivity are:
# 1. cp - Chest Pain type
# 2. thalach - maximum heart rate achieved
# 3. ca - number of major vessels coloured by flouroscopy
# 4. oldpeak - ST depression induced by exercise relative to rest
```

```
[123]: # Run the model on the test data to check the accuracy of the predictions
```

```
pred_array = model6.predict(x_test)
pred_array
```

```
[123]: array([0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0,
        1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0,
        1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 1])
```

```
[124]: # Create a dataframe to combine the predictions with the test data
```

```
pd.DataFrame(pred_array, columns = ['Pred_CVD'])
```

```
[124]:      Pred_CVD
0         0
1         1
2         0
3         0
4         0
..        ...
56        1
57        0
58        1
59        0
60        1
```

[61 rows x 1 columns]

```
[125]: # Create the new test data with the predicted values column

test_final = pd.concat([df_test.reset_index(drop = True),
                        pd.DataFrame(pred_array, columns = ['Pred_CVD'])], axis=
↪= 1)
test_final
```

```
[125]:      age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  \
0      62    0   0      160    164    0         0      145     0     6.2
1      56    1   1      130    221    0         0      163     0     0.0
2      48    1   0      130    256    1         0      150     1     0.0
3      60    1   0      130    206    0         0      132     1     2.4
4      50    1   0      150    243    0         0      128     0     2.6
..    ...  ...  ..      ...    ...    ...      ...    ...
56     63    1   3      145    233    1         0      150     0     2.3
57     57    1   0      110    335    0         1      143     1     3.0
58     38    1   3      120    231    0         1      182     1     3.8
59     57    1   0      150    276    0         0      112     1     0.6
60     64    0   0      180    325    0         1      154     1     0.0
```

```
      slope  ca  thal  target  Pred_CVD
0         0   3    3        0         0
1         2   0    3        1         1
2         2   2    3        0         0
3         1   2    3        0         0
4         1   0    3        0         0
..    ...  ...  ...      ...      ...
56         0   0    1        1         1
57         1   1    3        0         0
58         1   0    3        0         1
59         1   1    1        0         0
60         2   0    2        1         1
```


[61 rows x 15 columns]

```
[126]: # The predictions and the target variables are binary
# Will therefore create a new Boolean column with True and False
# This will help to calculate the accuracy of the predictions

test_final['Predicted'] = np.where((test_final.target == test_final.Pred_CVD),
    ↪ 'True', 'False')
test_final
```

```
[126]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	62	0	0	160	164	0	0	145	0	6.2	
1	56	1	1	130	221	0	0	163	0	0.0	
2	48	1	0	130	256	1	0	150	1	0.0	
3	60	1	0	130	206	0	0	132	1	2.4	
4	50	1	0	150	243	0	0	128	0	2.6	
..	
56	63	1	3	145	233	1	0	150	0	2.3	
57	57	1	0	110	335	0	1	143	1	3.0	
58	38	1	3	120	231	0	1	182	1	3.8	
59	57	1	0	150	276	0	0	112	1	0.6	
60	64	0	0	180	325	0	1	154	1	0.0	

	slope	ca	thal	target	Pred_CVD	Predicted
0	0	3	3	0	0	True
1	2	0	3	1	1	True
2	2	2	3	0	0	True
3	1	2	3	0	0	True
4	1	0	3	0	0	True
..
56	0	0	1	1	1	True
57	1	1	3	0	0	True
58	1	0	3	0	1	False
59	1	1	1	0	0	True
60	2	0	2	1	1	True

[61 rows x 16 columns]

```
[127]: # Total the True and False in the predicted values

test_final.Predicted.value_counts()

# There are 49 True (which is 21 (TP) + 28 (TN) as per confusion matrix)
# The false predictions are 12 (which is 9 (FP) and 3 (FN) as per confusion
    ↪ matrix)
```

```
[127]: True      49
      False    12
      Name: Predicted, dtype: int64
```

```
[128]: # Calculate the error False/Total

      12/61

      # Error in prediction is 19.67%
```

```
[128]: 0.19672131147540983
```

```
[129]: # Accuracy of prediction is 1 - Error

      1 - 12/61

      # Accuracy of prediction is 80.32%
      # this matches with the train data accuracy of 80.32%
```

```
[129]: 0.8032786885245902
```

```
[130]: # Our model is therefore accurate in prediction of the heart condition and can
      ↪ be used
```

```
[131]: # Check the Coefficient of Determination - R-Squared value of the model on the
      ↪ test data

      from sklearn.metrics import r2_score
```

```
[132]: # Coefficient of Determination - R-Squared

      r2_score(test_final.target, test_final.Pred_CVD)
```

```
[132]: 0.2129032258064516
```

```
[133]: # Since the number of variables are 13, we would need to arrive at the adjusted
      ↪ R-Squared

      1 - r2_score(test_final.target, test_final.Pred_CVD)
```

```
[133]: 0.7870967741935484
```

```
[134]: # The Adjusted R_Squared value is 78.71%
      # This value is very close to the accuracy of the model
      # The model can thus be deployed for accurate predictions of heart disease
```

```
[135]: # Final Conclusions

# The top 4 significant features (out of the 13) that helped predict a positive/
↳negative diagnosis are:
# 1. cp - chest pain type
# 2. thalach - maximum heart rate achieved
# 3. ca - number of major vessels and
# 4. oldpeak - ST depression induced by exercise relative to rest

# Random Forest algorithm yields the maximum accuracy of 80.32%
# The model worked well on the test data and thus assumed to be accurate for use
# Coefficient of Determination - R-Squared is 0.213
# Adjusted R-Squared is 78.70
# The model is thus well suited to be deployed in predicting heart conditions
```